


Predictive Maintenance System for Engine Condition Monitoring



Prepared & submitted by
Gopikrishna Rajendran
31-May-2026

Capstone Project – PGP AIML
Interim Submission Report (31-May-2026)

Table of Contents

Topic	Sub Topic
1. Project Overview	-
2. Business Problem	-
3. Objective	-
4. Data Description:	-
5. Data Registration	5.1 Folder Structure
	5.2 Registering Dataset in Hugging Face Space
6. Exploratory Data Analysis (EDA)	6.1 Dataset Overview
	6.2 Target Variable Distribution
	6.3 Feature Distribution
	6.4 Correlation Analysis
	6.5 Feature vs Target Analysis
	6.6 EDA Summary
7. Data Preparation	7.1 Data Overview
	7.2 Data Quality Validation
	7.3 Feature Verification
	7.4 Outlier Assessment
	7.5 Feature Selection
	7.6 Train-Test Split
	7.7 Data Preparation Summary
8. Model Building	8.1 Overview
	8.2 Logistic Regression
	8.3 Decision Tree
	8.4 Random Forest
	8.5 AdaBoost
	8.6 Gradient Boosting
	8.7 XGBoost
	8.8 Model Building Summary
9. Model Evaluation and Comparison	9.1 Evaluation Metrics
	9.2 Model Performance Comparison
	9.3 Comparative Analysis
	9.4 Model Selection
	9.5 Evaluation Summary
10. Hyperparameter Tuning	10.1 Overview
	10.2 Tuning Methodology
	10.3 Parameters Evaluated
	10.4 Best Parameters
	10.5 Hyperparameter Tuning Results
	10.6 Hyperparameter Tuning Summary
11. Conclusion:	-
12. Business Report Quality Checklist	-

1. Project Overview

This project focuses on developing a machine learning-based predictive maintenance solution for engine condition monitoring. Using historical engine sensor data, multiple machine learning models were developed and evaluated to predict potential engine failures. The selected model was registered and deployed as a live web application, enabling real-time engine condition assessment and supporting proactive maintenance decisions.

2. Business Problem

Unexpected engine failures can result in costly repairs, operational downtime, reduced productivity, and safety concerns. Traditional maintenance approaches often rely on fixed schedules or reactive actions after a failure occurs, which may lead to unnecessary maintenance costs or unplanned breakdowns. Organizations require a predictive maintenance solution that can leverage engine sensor data to identify potential failures in advance. By analyzing key engine parameters such as RPM, pressure, and temperature, machine learning techniques can help predict engine condition and support proactive maintenance decisions.

3. Objective

The objective of this project is to develop a machine learning-based predictive maintenance system capable of predicting engine failures using operational and sensor data. Multiple machine learning algorithms will be evaluated and compared to identify the most suitable model for failure prediction.

The deployed solution aims to:

- Reduce unexpected engine breakdowns and maintenance costs.
- Improve equipment reliability and operational efficiency.
- Support proactive maintenance planning.
- Enable data-driven decision-making through real-time engine condition assessment.


4. Data Description:

- Engine_RPM: The number of revolutions per minute (RPM) of the engine, indicating engine speed. It is defined in Revolutions per Minute (RPM).
- Lub_Oil_Pressure: The pressure of the lubricating oil in the engine, essential for reducing friction and wear. It is defined in bar or kilopascals (kPa)
- Fuel_Pressure: The pressure at which fuel is supplied to the engine, critical for proper combustion. It is defined in bar or kilopascals (kPa)
- Coolant_Pressure: The pressure of the engine coolant, affecting engine temperature regulation. It is defined in bar or kilopascals (kPa)
- Lub_Oil_Temperature: The temperature of the lubricating oil, which impacts viscosity and engine performance. It is defined in degrees Celsius (°C)
- Coolant_Temperature: The temperature of the engine coolant, crucial for preventing overheating. It is defined in degrees Celsius (°C)
- Engine_Condition: A categorical or numerical label representing the health of the engine, potentially indicating normal operation or various levels of wear and failure risks. It is defined as a categorical variable (0/1) representing a state such as "0 = Off/False/Active" and "1 = On/True/Faulty"

5. Data Registration

5.1 Folder Structure

- The project workspace was organized in PyCharm using a master folder (**predictive-maintenance-capstone**) with dedicated subfolders for data, notebooks, source code, deployment artifacts, reports, and model files to ensure a structured development workflow



5.2 Registering Dataset in Hugging Face Space

- The dataset was published and versioned in a Hugging Face Dataset Repository to enable centralized access, reproducibility, and future integration with deployment workflows.
- Dataset Repo Link: <https://huggingface.co/datasets/gkrishna4346/predictive-maintenance-dataset>

Attribute	Value
Dataset Name	engine_data.csv
Records	19535
Features	7
Repository Type	Hugging Face Dataset
Access	Private/Public (whichever applies)

Screenshot of Hugging Face Space:

A screenshot of the Hugging Face dataset viewer for the 'predictive-maintenance-dataset' by user 'gkrishna4346'. The interface shows the dataset card with a search bar, navigation tabs (Tasks, Modalities, Formats, Languages, Size, Libraries, License), and a 'Dataset card' tab selected. The 'Dataset Viewer' section displays a table of data rows with columns: Engine xpr (int64), Lub oil pressure (float64), Fuel pressure (float64), Coolant pressure (float64), Lub oil temp (float64), Coolant temp (float64), and Engine Condition (int64). The table shows 15.6k rows in the 'train' split. On the right, there are buttons for 'Copy to bucket', 'Use this dataset', 'Edit dataset card', and a 'Downloads last month' counter showing 20. Below that, a summary box indicates 'Number of rows: 19,535' and 'Total file size: 2.61 MB'.

6. Exploratory Data Analysis (EDA)


The exploratory analysis helped identify key patterns in engine operating parameters and their relationship with engine condition. The dataset showed no major data quality issues, and feature analysis revealed meaningful variations between normal and failure cases, supporting the use of machine learning for predictive maintenance.

6.1 Dataset Overview

Metric	Value
Dataset Name	engine_data.csv
Total Records	19,535
Total Features	7
Feature Types	All Numerical
Input Features	6
Target Variable	1
Missing Values	None (Dataset is complete)
Duplicate Records	None (All records are unique)

- **Independent Variables:** Engine RPM, Lub Oil Pressure, Fuel Pressure, Coolant Pressure, Lub Oil Temperature, Coolant Temperature & the data type is Numerical (Continuous) for all
- **Target Variable:** Engine Condition & the data type is Binary (0/1)

6.2 Target Variable Distribution



Key Observations:

- Moderate class imbalance exists (63% vs 37%) — not severe but warrants attention during model evaluation
- Class 0 (Normal Operating Condition) is the minority class with 7,218 records
- Class 1 (Requires Maintenance) is the majority class with 12,317 records

6.3 Feature Distribution

Feature	Distribution Shape	Outlier Presence
Engine RPM	Positively Skewed	Yes — Extreme high values (61–2239 RPM)
Lub Oil Pressure	Positively Skewed	Yes — Upper end outliers
Fuel Pressure	Positively Skewed	Yes — Multiple upper outliers
Coolant Pressure	Positively Skewed	Yes — Occasional high readings
Lub Oil Temperature	Concentrated range	Yes — Clustered between 82–90
Coolant Temperature	Moderate spread	Yes — Few extreme outliers

Key observations:


- All features are continuous and numerical; no categorical encoding required.

6.4 Correlation Analysis

A Pearson correlation analysis was conducted to measure linear relationships between features and the target variable. The analysis indicated low multicollinearity among features, enabling retention of all variables for model development.

Feature	Correlation with Target Variable	Strength
Engine RPM	-0.27	Moderate Negative
Fuel Pressure	0.12	Weak Positive
Lub Oil Pressure	0.06	Negligible
Lub Oil Temperature	-0.09	Negligible
Coolant Temperature	-0.05	Negligible
Coolant Pressure	-0.02	Negligible

Supporting Chart:



Key Observations:


- Engine RPM is the strongest linear predictor (correlation = -0.27), showing that higher RPM is associated with the maintenance-required condition
- Fuel Pressure has a weak but positive correlation (+0.12) with the normal operating condition
- No pair of features shows high inter-feature correlation, confirming low multicollinearity
- Weak correlations suggest that non-linear models may be more effective in capturing feature-target relationships

6.5 Feature vs Target Analysis

Bivariate analysis (Feature vs Engine Condition) was conducted using boxplots to examine how each feature varies across the two engine condition classes.

Feature	Class 0 (Normal Operating Condition)	Class 1 (Requires Maintenance)
Engine RPM	≈ 885	≈ 736
Fuel Pressure	≈ 6.24	≈ 6.90
Lub Oil Pressure	Lower	Slightly Higher
Coolant Pressure	Similar	Similar
Lub Oil Temperature	Similar	Similar
Coolant Temperature	Similar	Similar

Supporting Chart:



Note: Charts added only for the Comparison of Feature vs Engine Conditions (across classes) which has significant visibility

Key Observations:

- Engine RPM exhibited the most visible separation between classes
- Fuel Pressure showed meaningful variation across engine conditions
- Lub Oil Pressure shows a slight increase in Class 1, though class overlap is substantial.
- Rest all features exhibit nearly identical distributions across both classes.
- Engine failure prediction depends on the combined influence of multiple sensor measurements
- Multivariate Pairplot analysis confirmed that no strong clustering exists — ensemble or non-linear models are preferred

6.6 EDA Summary

Aspect	Finding
Dataset Size	19,535 records × 7 features (6 predictor variables + 1 target variable)
Data Quality	No missing values or duplicate records; dataset is clean and ready for analysis
Target Variable	Engine Condition (Binary: 0 = Normal Operating Condition, 1 = Maintenance Required)
Class Distribution	Normal Operating Condition: 36.95% (7,218 records) vs Maintenance Required: 63.05% (12,317 records); moderate class imbalance observed
Feature Types	All features are numerical: Engine RPM, Lub Oil Pressure, Fuel Pressure, Coolant Pressure, Lub Oil Temperature, and Coolant Temperature
Key Predictors	Engine RPM showed the strongest relationship with the target variable (corr = -0.27), followed by Fuel Pressure (corr = +0.12)
Multicollinearity	Low multicollinearity observed; no strong inter-feature correlations detected, therefore all features were retained
Outliers	Outliers were identified in multiple features; retained because tree-based models are generally robust to outlier effects
Skewness	Most numerical features exhibited positive skewness, indicating non-normal distributions
ML Readiness	Dataset is suitable for classification tasks. Ensemble techniques such as Random Forest, AdaBoost, and XGBoost are recommended due to their ability to capture complex relationships and improve predictive performance
Next Steps	Develop and evaluate classification models (Decision Tree, Random Forest, AdaBoost, and XGBoost) using Precision, Recall, F1-Score, ROC-AUC, and Confusion Matrix metrics to identify the best-performing predictive maintenance solution

Aspect	Finding
Additional Observation	<i>The dataset contains a higher proportion of maintenance-required instances compared to normal operating conditions. While this distribution may differ from typical real-world operating environments, all analyses and model development activities were performed using the class definitions and distribution provided in the dataset</i>

7. Data Preparation

7.1 Data Overview

Data preparation was performed to ensure that the dataset was suitable for machine learning model development. The process included data quality checks, validation of feature types, assessment of missing values and duplicates, outlier analysis, feature selection, and train-test splitting.

7.2 Data Quality Validation

The dataset was examined to identify potential data quality issues that could impact model performance.

Activities Performed:

- Verified dataset dimensions and structure
- Reviewed feature data types
- Checked for missing values
- Checked for duplicate records
- Validated target variable consistency

Findings:

Validation Check	Result
Missing Values	None
Duplicate Records	None
Invalid Data Types	None
Target Variable Issues	None

7.3 Feature Verification

- All six predictor variables were numerical and represented engine operating parameters.
- No feature encoding was required because all predictor variables were already numerical.

7.4 Outlier Assessment

Boxplot analysis performed during exploratory data analysis identified the presence of outliers in multiple numerical features.

Findings:

- Engine RPM contained several high-value observations.
- Fuel Pressure and Lub Oil Pressure showed upper-end outliers.
- Temperature-related features contained a few extreme values.

Note:

Since the selected machine learning algorithms include Decision Tree, Random Forest, AdaBoost, and XGBoost, which are generally robust to outliers, no outlier treatment was applied. All records were retained to preserve operational variability.

7.5 Feature Selection

Correlation analysis and exploratory data analysis were used to evaluate feature relevance.

Findings:

- Engine RPM showed the strongest relationship with the target variable.
- Fuel Pressure exhibited a weak positive relationship with engine condition.
- No strong multicollinearity was observed among features.

Therefore, all six predictor variables were retained for model development.

7.6 Train-Test Split

To evaluate model performance on unseen data, the dataset was divided into training and testing subsets.

Parameter	Value
Training Data	80%
Testing Data	20%
Random State	42

The training dataset was used to build machine learning models, while the testing dataset was reserved for model evaluation.

7.7 Data Preparation Summary

Step	Outcome
Missing Value Check	No missing values identified
Duplicate Check	No duplicate records found
Feature Verification	All variables numerical
Outlier Assessment	Outliers retained
Feature Selection	All six predictors retained
Train-Test Split	80:20
Dataset Readiness	Suitable for model development

The dataset required minimal preprocessing due to its high quality and complete structure. Following validation and preparation activities, the dataset was considered ready for machine learning model development and evaluation.

8. Model Building

8.1 Overview

Multiple machine learning classification algorithms were developed and trained to predict engine condition based on operational sensor measurements. Both single-tree and ensemble learning techniques were evaluated to identify the most effective predictive maintenance solution.

The selected models include:

- Logistic Regression
- Decision Tree
- Random Forest
- AdaBoost
- Gradient Boosting
- XGBoost

These algorithms were chosen because they can effectively handle numerical data, capture non-linear relationships, and provide strong classification performance.

8.2 Logistic Regression

Model Rationale: Logistic Regression was selected as the baseline statistical classification model due to its simplicity, interpretability, and effectiveness in binary classification problems.

Key Characteristics:

- Suitable for binary classification
- Provides interpretable predictions
- Establishes a benchmark for evaluating advanced models.

8.3 Decision Tree

Model Rationale: A Decision Tree classifier was developed to capture non-linear relationships between engine operating parameters and engine condition.

Key Characteristics:

- Easy to interpret
- Handles numerical features without scaling
- Serves as a baseline tree-based model

8.4 Random Forest

Model Rationale: Random Forest was developed to improve predictive performance by combining multiple decision trees using a bagging approach.

Key Characteristics:

- Reduces overfitting
- Improves prediction stability
- Captures complex feature interactions

8.5 AdaBoost

Model Rationale: AdaBoost was developed as a boosting-based ensemble model that focuses on correcting previously misclassified observations.

Key Characteristics:

- Sequential learning approach
- Improves classification performance
- Effective for structured tabular datasets

8.6 Gradient Boosting

Model Rationale: Gradient Boosting was developed to build a strong predictive model by sequentially minimizing prediction errors.

Key Characteristics:

- Learns from residual errors
- Captures non-linear relationships
- Generally provides strong classification performance

8.7 XGBoost

Model Rationale: XGBoost was developed as an advanced gradient boosting framework designed for efficiency, scalability, and predictive accuracy.

Key Characteristics:

- Regularization reduces overfitting
- Handles complex feature interactions
- Widely used in machine learning competitions and industry applications

8.8 Model Building Summary

Model	Category	Type
Logistic Regression	Linear Model	Baseline
Decision Tree	Tree-Based	Baseline Tree Model
Random Forest	Ensemble	Bagging
AdaBoost	Ensemble	Boosting
Gradient Boosting	Ensemble	Boosting
XGBoost	Ensemble	Advanced Boosting

Key Observation:

The selected models represent a progression from a simple linear classifier to advanced ensemble learning techniques. This approach enables objective comparison across different algorithm families and supports the identification of the most effective predictive maintenance solution.

9. Model Evaluation and Comparison

9.1 Evaluation Metrics

To evaluate model performance, multiple classification metrics were considered. Since the dataset exhibits moderate class imbalance, relying solely on accuracy may not provide a complete assessment of model effectiveness.

The following metrics were used:

Metric	Description
Accuracy	Overall proportion of correctly classified observations
Precision	Ability to correctly identify maintenance-required engines
Recall	Ability to detect actual maintenance-required engines
F1-Score	Harmonic mean of Precision and Recall
ROC-AUC	Ability to distinguish between the two classes across different thresholds

Metrics such as **Precision, Recall, F1-Score, and ROC-AUC** were given greater importance than accuracy because they provide a more comprehensive assessment of classification performance under class imbalance conditions.

9.2 Model Performance Comparison

Model	Accuracy	ROC-AUC	Precision	Recall	F1-Score
Logistic Regression	0.6560	0.6804	0.6737	0.8810	0.7635
Decision Tree	0.5930	0.5610	0.6750	0.6837	0.6793
Random Forest	0.6453	0.6773	0.6887	0.7978	0.7393
AdaBoost	0.6688	0.7050	0.6979	0.8368	0.7611
Gradient Boosting	0.6647	0.7035	0.6920	0.8437	0.7603
XGBoost	0.6455	0.6700	0.6925	0.7873	0.7368

9.3 Comparative Analysis

Logistic Regression

- Achieved the highest Recall (88.10%).
- Achieved the highest F1-Score (76.35%).
- Demonstrated strong ability to identify maintenance-required engines.
- Performance suggests that the relationship between features and target variables can be partially captured using a linear decision boundary.

Decision Tree

- Produced the lowest Accuracy and ROC-AUC among all models.
- More susceptible to overfitting compared to ensemble methods.
- Served as a useful baseline tree-based model.

Random Forest

- Improved performance over Decision Tree.
- Demonstrated better stability through ensemble learning.
- Produced balanced classification performance across evaluation metrics.

AdaBoost

- Achieved the highest Accuracy (66.88%).
- Achieved the highest ROC-AUC (70.50%).
- Achieved the highest Precision (69.79%).
- Demonstrated strong overall performance across all evaluation metrics.

Gradient Boosting

- Delivered performance very close to AdaBoost.
- Achieved strong Recall and ROC-AUC values.
- Demonstrated effective learning of complex relationships in the dataset.

XGBoost

- Produced competitive results but did not outperform AdaBoost or Gradient Boosting.
- Demonstrated reasonable classification capability and model stability.

9.4 Model Selection

Although Logistic Regression achieved the highest Recall and F1-Score, AdaBoost was selected as the final model because it provided the best overall balance across evaluation metrics.

Reasons for Selecting AdaBoost:

Criterion	Observation
Accuracy	Highest among all models
ROC-AUC	Highest among all models
Precision	Highest among all models
Recall	Strong performance
Model Stability	Consistent across metrics
Interpretability	Easier deployment compared to more complex boosting frameworks

AdaBoost demonstrated superior discriminative capability while maintaining a balanced trade-off between false positives and false negatives, making it the most suitable model for the predictive maintenance use case.

9.5 Evaluation Summary

The comparative analysis showed that ensemble learning techniques consistently outperformed the standalone Decision Tree model. Among the evaluated algorithms, **AdaBoost** achieved the strongest overall performance, delivering the highest Accuracy, ROC-AUC, and Precision while maintaining competitive Recall and F1-Score values. Based on these results, **AdaBoost** was selected as the **final model** for subsequent deployment and predictive maintenance analysis.

10. Hyperparameter Tuning

10.1 Overview

After evaluating multiple machine learning algorithms, AdaBoost emerged as the best-performing model based on Accuracy, Precision, and ROC-AUC. To further optimize its predictive performance, hyperparameter tuning was performed using Grid Search Cross Validation (GridSearchCV).

The objective was to identify the optimal combination of model parameters that maximized classification performance while maintaining generalization capability.

10.2 Tuning Methodology

GridSearchCV was employed to systematically evaluate multiple parameter combinations for the AdaBoost classifier.

Benefits of Grid Search

- Automated parameter optimization
- Systematic evaluation of parameter combinations
- Cross-validation-based performance assessment
- Improved model reliability and generalization

The ROC-AUC score was selected as the optimization metric because it provides a robust measure of the model's ability to distinguish between maintenance-required and normal operating conditions.

10.3 Parameters Evaluated

The following AdaBoost hyperparameters were explored:

Parameter	Description
n_estimators	Number of weak learners used during boosting
learning_rate	Controls the contribution of each weak learner

Candidate Values:

Parameter	Values Evaluated
n_estimators	50, 100, 150, 200
learning_rate	0.01, 0.1, 0.5, 1.0

10.4 Best Parameters

The Grid Search process identified the following optimal parameter combination:

Parameter	Best Value
learning_rate	0.1
n_estimators	200

Best Cross-Validation Performance:

Metric	Value
ROC-AUC Score	0.7012

10.5 Hyperparameter Tuning Results

The tuning process successfully identified an improved AdaBoost configuration that achieved strong classification performance across validation folds.

Key Observations:

- Increasing the number of estimators improved the model's learning capability.
- A moderate learning rate of 0.1 provided a good balance between learning speed and model stability.
- The tuned AdaBoost model achieved a cross-validated ROC-AUC score of 0.7012.
- The optimized model demonstrated strong capability in distinguishing maintenance-required engines from normal operating engines.

10.6 Hyperparameter Tuning Summary

Activity	Outcome
Model Tuned	AdaBoost
Optimization Technique	GridSearchCV
Validation Method	Cross Validation
Best Learning Rate	0.1
Best Number of Estimators	200
Best ROC-AUC Score	0.7012
Final Model Status	Selected for Deployment

11. Conclusion:

Hyperparameter tuning enhanced the AdaBoost model by identifying the optimal parameter configuration of **learning_rate = 0.1** and **n_estimators = 200**. The tuned model achieved a ROC-AUC score of **0.7012** and was selected as the final predictive maintenance model for deployment.

Figure 10.1: Optimal AdaBoost Hyperparameters Identified using GridSearchCV

12.3. Best Parameters

 # Best Parameters

```
print("Best Parameters:")
print(grid_search.best_params_)

print("\nBest ROC-AUC Score:")
print(round(grid_search.best_score_, 4))
```

... Best Parameters:
{'learning_rate': 0.1, 'n_estimators': 200}

Best ROC-AUC Score:
0.7012

[+ Code](#) [+ Text](#)

12. Business Report Quality Checklist

Topic	Checklist	Status
1. Report Structure	Cover Page included	Yes
	Table of Contents included	Yes
	Page numbers present	Yes
	Section numbering is sequential	Yes
	Headings follow consistent formatting	Yes
	Report does not exceed page limit	Yes
	Figures and tables are numbered	Yes
	Figure/Table captions provided	Yes
2. Business Understanding	Business problem clearly defined	Yes
	Business impact explained	Yes
	Project objective clearly stated	Yes
	Success criteria explained	Yes
	Predictive maintenance use case justified	Yes
	Business value communicated	Yes
3. Data Understanding	Dataset description provided	Yes
	Data source identified	Yes
	Target variable explained	Yes
	Feature descriptions included	Yes
	Dataset size documented	Yes
	Assumptions disclosed	Yes
	Data limitations discussed	Yes
4. Exploratory Data Analysis (EDA)	Dataset overview included	Yes
	Target distribution visualized	Yes
	Feature distributions analyzed	Yes
	Correlation analysis performed	Yes
	Feature vs Target analysis included	Yes
	EDA observations documented	Yes
	EDA summary table included	Yes
5. Data Preparation	Missing value analysis performed	Yes
	Duplicate analysis performed	Yes
	Outlier assessment discussed	Yes
	Feature selection justified	Yes
	Train-test split documented	Yes
	Data readiness summary provided	Yes
6. Model Building	Multiple models evaluated	Yes
	Model rationale explained	Yes
	Model building process documented	Yes
	Algorithms clearly identified	Yes
	Hyperparameter tuning included	Yes
	Best parameters reported	Yes
	Experimentation process explained	Yes
7. Model Evaluation	Accuracy reported	Yes
	Precision reported	Yes
	Recall reported	Yes
	F1-Score reported	Yes
	ROC-AUC reported	Yes
	Comparison table included	Yes
	Best model selection justified	Yes
	Metrics interpreted correctly	Yes